

Cloud-Based College Campus Management System

A Serverless Multi-Tier Architecture on Amazon
Web Services

Submitted by:

Karthik Chakala

Roll Number: 22CSE0104

Course:

Cloud Computing (CS401)

Department of Computer Science & Engineering

National Institute of Technology

Academic Year 2025–2026

Abstract

Higher education institutions face significant operational challenges when managing campus administrative operations, student enrollments, course schedules, daily attendance tracking, assignment evaluation, and examination publishing on traditional, legacy on-premises server infrastructure. Static server hardware suffers from high capital expenditure, single points of failure, administrative maintenance overhead, and resource under-utilization during off-peak periods, paired with database connection exhaustion during peak registration windows.

This report presents the design, implementation, cloud deployment, and verification of the **College Campus Management System**, an enterprise-grade, full-stack, cloud-native application deployed on Amazon Web Services (AWS). To eliminate continuous idle compute billing and server maintenance overhead, the cloud architecture deliberately excludes Amazon EC2 in favor of a serverless, managed multi-tier design. The system pairs a React 18 Single Page Application hosted via Amazon S3 Static Website Hosting with a serverless backend comprising Amazon API Gateway REST APIs, AWS Lambda microservices (Node.js 20.x, Express, Prisma ORM), and an Amazon RDS for PostgreSQL 15.3 relational database isolated inside private Amazon VPC subnets. Identity authentication is managed via Amazon Cognito User Pools issuing RSA-256 JWT tokens, while event alerts are dispatched asynchronously via Amazon Simple Notification Service (Amazon SNS). Telemetry logs and performance metrics are centrally monitored using Amazon CloudWatch, bounded by AWS Identity and Access Management (AWS IAM) least-privilege security policies.

Comprehensive system verification confirmed the functionality of all 44 REST API routes, 17 database relational models, and 22 automated integration tests with 100% pass rates across Student, Faculty, and Administrator user roles. Quantitative financial modeling demonstrated a 74% compute cost reduction compared to conventional virtual machine deployments, enabling the application prototype to operate completely free of charge under the AWS Free Tier. The report provides a deep technical analysis of each AWS primitive, details inter-service workflows, documents empirical implementation trade-offs, and validates how serverless cloud architecture effectively supports modern institutional enterprise software.

Contents

Abstract	i
1 Introduction	1
1.1 Background and Context	1
1.2 Problem Statement	1
1.3 Project Motivation	2
1.4 Project Scope	2
1.5 Report Organization	3
2 Cloud Computing Fundamentals and Rationale	4
2.1 Service Models in Cloud Computing	4
2.2 Core Attributes of Cloud Architecture	4
2.3 Application to Institutional Campus Management	5
3 AWS Selection Rationale and Ecosystem	6
3.1 Why Amazon Web Services (AWS)?	6
3.2 AWS Regional Architecture	6
4 Deep Dive into Utilized AWS Services	7
4.1 Amazon Simple Storage Service (Amazon S3)	7
4.2 Amazon Relational Database Service (Amazon RDS for PostgreSQL)	8
4.3 AWS Lambda	9
4.4 Amazon API Gateway	9
4.5 Amazon Cognito	10
4.6 Amazon Simple Notification Service (Amazon SNS)	11
4.7 Amazon CloudWatch	11
4.8 AWS Identity and Access Management (AWS IAM)	12
4.9 Amazon Virtual Private Cloud (Amazon VPC)	13
4.10 Why Amazon EC2 Was Not Selected	13
5 System Requirements Analysis	15
5.1 Functional Requirements	15
5.1.1 Student Role Requirements	15
5.1.2 Faculty Role Requirements	15
5.1.3 Administrator Role Requirements	16
5.2 Non-Functional Requirements	16

6	System Architecture and Topology	17
6.1	Overall Architecture Overview	17
6.2	Component Tier Breakdown	18
6.2.1	Presentation Tier (Frontend)	18
6.2.2	API Gateway and Compute Tier (Backend)	18
6.2.3	Persistence Tier (Database and Object Storage)	18
6.2.4	Authentication and Messaging Tier	18
6.2.5	Security and Management Tier	18
6.3	AWS Service Communication Topology	18
7	Database Schema and Service Implementation	20
7.1	PostgreSQL Relational Schema Design	20
7.2	Service Implementation Evidence	22
7.2.1	Amazon S3 Static Website & Document Storage	22
7.2.2	Amazon RDS for PostgreSQL	22
7.2.3	AWS Lambda Serverless Compute	22
7.2.4	Amazon API Gateway REST API	22
7.2.5	Amazon Cognito Identity Authentication	22
8	Service Integration and Multi-Service Workflows	25
8.1	Workflow 1: Student Authentication and Session Acquisition	25
8.2	Workflow 2: Student Assignment File Upload and Submission	26
8.3	Workflow 3: Faculty Grading and SNS Notification Dispatch	26
8.4	Workflow 4: Administrative Reporting and CSV Export	27
9	Cloud Security Architecture and Shared Responsibility	28
9.1	AWS Shared Responsibility Model	28
9.2	Identity Security and Access Control	28
9.2.1	IAM Least-Privilege Execution Roles	28
9.2.2	Role-Based Access Control (RBAC)	28
9.3	Network Isolation and Database Security	29
9.4	Data Protection and Encryption	29
10	Quantitative Cost Analysis and Economic Model	30
10.1	Serverless vs. Provisioned Economic Model	30
10.2	Detailed AWS Pricing Breakdown	30
10.3	Cost Advantage of Excluding Amazon EC2	31
11	Cloud Deployment Procedure	32
11.1	Step-by-Step Deployment Pipeline	32
11.1.1	Phase 1: VPC and Relational Database Provisioning	32
11.1.2	Phase 2: Database Schema Migration and Seeding	32
11.1.3	Phase 3: AWS Lambda Backend Packaging and Deployment	32
11.1.4	Phase 4: API Gateway Integration	33
11.1.5	Phase 5: S3 Static Frontend Deployment	33

12 Technical Observations and Engineering Trade-Offs	34
12.1 Observed Infrastructure Behaviors	34
12.2 Inferred Architectural Patterns	34
12.3 Expected Theoretical Models	35
13 Testing and Verification Validation	36
13.1 Automated Integration Test Suite	36
13.2 Database Constraint Verification	36
13.3 End-to-End User Workflow Execution	37
14 Results and Discussion	38
14.1 System Performance and Latency Metrics	38
14.2 Comparative Evaluation: Local vs. AWS Cloud Architecture	38
15 Limitations and Future Enhancements	40
15.1 System Limitations	40
15.2 Future Enhancements	40
16 Conclusion	42
16.1 Summary of Achievements	42
16.2 Closing Remarks	42

List of Figures

6.1	Overall Cloud Architecture Topology for College Campus Management System	17
6.2	AWS Ecosystem Inter-Service Communication and Security Topology . .	19
7.1	Amazon S3 Static Website Hosting Management Console View	20
7.2	Amazon RDS for PostgreSQL Database Instance Configuration Console View	22
7.3	AWS Lambda Backend Function Configuration Console View	23
7.4	Amazon API Gateway REST API Routes and Integration Console View	23
7.5	Amazon Cognito User Pool Configuration Console View	24
11.1	CloudWatch Monitoring Dashboard for AWS Lambda and RDS Services Console View	33

List of Tables

2.1	Traditional On-Premises Deployment vs. Cloud-Native AWS Architecture	5
7.1	Core Database Schema Entities and Relational Constraints	21
10.1	AWS Service Pricing Breakdown and Monthly Cost Model	30
13.1	Automated Vitest Integration Test Suite Summary	36
14.1	Comparative Matrix: Local Monolithic Deployment vs. AWS Serverless Cloud Architecture	39

Chapter 1

Introduction

1.1 Background and Context

Higher education institutions rely heavily on centralized information systems to manage administrative operations, academic schedules, student enrollments, faculty evaluations, assignment grading, attendance monitoring, and examination records. Traditional campus management software was historically deployed on dedicated, physical on-premises servers located within institution data centers. While on-premises deployments offered perimeter-level physical control, they introduced severe operational challenges, including high capital expenditure (CapEx) for hardware procurement, single points of failure during hardware degradation, inflexible capacity limits during registration peak periods, and heavy maintenance overhead for system administrators.

Cloud computing has transformed institutional software engineering by replacing static physical hardware with elastic, programmable, and on-demand cloud infrastructure. Amazon Web Services (AWS) provides a comprehensive suite of cloud primitives—spanning storage, serverless compute, relational databases, identity management, notification routing, network isolation, and telemetry monitoring—enabling institutions to deploy high-availability, cost-effective campus management applications without purchasing or maintaining physical servers.

1.2 Problem Statement

Legacy institutional platforms struggle to maintain performance and consistency under fluctuating workload profiles. For instance, during enrollment registration windows or assignment submission deadlines, access requests surge by several orders of magnitude. Under an on-premises model, servers must be provisioned for peak load capacity, resulting in low average resource utilization (often below 15%) during normal operational periods. Conversely, under-provisioned infrastructure experiences server crashes, database connection exhaustions, slow response times, and unhandled requests during peak demand.

Furthermore, traditional web applications host monolithic backend services, database engines, file storage, and authentication on single virtual machines. This coupled architecture violates the principle of separation of concerns, exposes full server infrastructure to security threats, and makes disaster recovery and zero-downtime maintenance practically impossible.

1.3 Project Motivation

The motivation of this project is to design, implement, and validate a modern, full-stack, cloud-native **College Campus Management System** utilizing Amazon Web Services. The system is engineered to solve operational bottlenecks by leveraging cloud-native serverless and managed services:

- **Elasticity and Availability:** Decoupling static asset delivery from backend API handling and database persistence to guarantee system responsiveness under variable load.
- **Cost Optimization:** Adopting a serverless and managed cloud footprint (Amazon S3, AWS Lambda, Amazon API Gateway, Amazon RDS PostgreSQL) that eliminates the continuous idle running costs associated with always-on Amazon EC2 virtual machines.
- **Security and Role Isolation:** Enforcing robust authentication and role-based access control (RBAC) across three distinct institutional roles *Student*, *Faculty*, and *Administrator*.
- **Academic Workflow Automation:** Streamlining core institutional workflows including enrollment, daily attendance logging, assignment submission and grading, exam result publishing, broadcast announcements, and audit logging.

1.4 Project Scope

The scope of this project encompasses the complete lifecycle of a cloud-based campus management solution:

1. **Frontend Development:** Building a responsive Single Page Application (SPA) using React 18, TypeScript, Tailwind CSS, and Lucide icons, hosted via Amazon S3 Static Website Hosting.
2. **Backend Microservices API:** Developing 44 RESTful API endpoints using Node.js, Express, and Prisma ORM, deployed via AWS Lambda and exposed through Amazon API Gateway.
3. **Relational Database Architecture:** Designing a PostgreSQL database schema containing 17 models with foreign key cascades, unique indices, and transaction boundaries, managed by Amazon RDS for PostgreSQL.
4. **Cloud Service Integration:** Wiring Amazon Cognito for JWT authentication, Amazon SNS for notification publishing, Amazon CloudWatch for centralized logging, AWS IAM for least-privilege security, and Amazon VPC for private network isolation.
5. **System Verification:** Performing complete automated vitest integration testing, API route matrix validation, database constraint testing, and end-to-end user workflow execution.

1.5 Report Organization

The remainder of this report is structured as follows: Chapter 2 reviews core cloud computing paradigms and AWS design principles. Chapter 3 outlines the selection rationale for AWS. Chapter 4 provides an exhaustive 13-aspect technical analysis of all 9 utilized AWS services and explains why Amazon EC2 was excluded. Chapter 5 defines functional and non-functional system requirements. Chapter 6 presents the cloud system architecture and communication topology. Chapter ?? details the PostgreSQL database schema and ER diagram. Chapter 7 documents service-by-service implementation. Chapter 8 explains step-by-step inter-service workflows. Chapter ?? analyzes the serverless compute model. Chapter 9 covers cloud security and shared responsibility. Chapter 10 details quantitative cost analysis. Chapter 11 outlines deployment procedures. Chapter 12 presents empirical implementation observations. Chapter 13 details automated and E2E test results. Chapter 14 compares local versus cloud performance. Chapter 15 discusses system limitations and future enhancements. Chapter 16 concludes the report.

Chapter 2

Cloud Computing Fundamentals and Rationale

2.1 Service Models in Cloud Computing

Cloud computing abstracts raw physical infrastructure into accessible, programmable digital services. The Cloud Security Alliance (CSA) and NIST classify cloud offerings into three primary service models:

1. **Infrastructure as a Service (IaaS)**: Provides raw compute instances, virtual disks, and software-defined networking primitives. Users maintain full control over operating systems, runtime middleware, security patches, and application deployments, while the provider manages physical servers and data center facilities. (e.g., Amazon EC2, AWS EBS).
2. **Platform as a Service (PaaS)**: Delivers pre-configured execution environments where developers deploy code without managing OS maintenance or web server binaries. (e.g., AWS Elastic Beanstalk).
3. **Software as a Service (SaaS)**: Delivers fully managed, consumer-facing end-user applications hosted entirely in the cloud. (e.g., Google Workspace, Microsoft 365).

In recent years, **Serverless Computing** (Function-as-a-Service, FaaS) has emerged as an evolution of PaaS. In serverless computing, developers deploy stateless event-driven functions (e.g., AWS Lambda) while the cloud vendor manages automatic scaling, process isolation, OS patching, and execution scheduling.

2.2 Core Attributes of Cloud Architecture

Modern cloud applications leverage key architectural attributes that distinguish them from traditional deployments:

- **Scalability and Elasticity**: Scalability refers to the system's ability to handle growing workloads by increasing capacity. Elasticity is the dynamic adaptation of capacity scaling out during traffic surges and scaling in when load subsides ensuring resource allocation strictly tracks demand.

- **High Availability and Fault Tolerance:** Cloud infrastructure is distributed across geographically isolated Availability Zones (AZs). Applications designed with multi-AZ replication survive hardware failures without downtime.
- **Pay-As-You-Go Financial Model:** Replaces upfront capital expenditure (CapEx) with operational expenditure (OpEx), charging consumers only for exact compute duration, request volume, and bytes stored.
- **Shared Responsibility Model:** Defines the division of security obligations between the cloud provider and the consumer. The provider guarantees security *of* the cloud (hardware, facilities, hypervisors), while the consumer manages security *in* the cloud (IAM policies, database encryption, network access rules, application code).

2.3 Application to Institutional Campus Management

The College Campus Management System exemplifies an ideal workload for cloud migration:

Table 2.1: Traditional On-Premises Deployment vs. Cloud-Native AWS Architecture

Architectural Attribute	Traditional On-Premises Setup	Cloud-Native AWS Architecture
Capacity Planning	Over-provisioned for peak load; idle 85% of time.	Elastic serverless scaling tracking live HTTP request rate.
Capital Cost (CapEx)	High upfront server procurement and rack costs.	\$0 upfront; operational billing based on usage.
High Availability	Single server failure causes institutional downtime.	Multi-AZ replication for RDS, S3, and Lambda functions.
Asset Storage	Local hard drives prone to disk corruption.	Amazon S3 with 99.999999999% (11 9s) durability.
Security Maintenance	Manual OS security patching and firewall rules.	Automated AWS managed OS patching and IAM RBAC rules.

Chapter 3

AWS Selection Rationale and Ecosystem

3.1 Why Amazon Web Services (AWS)?

Amazon Web Services (AWS) was selected as the cloud infrastructure provider for this system due to its comprehensive service ecosystem, industry-leading reliability, fine-grained identity security, and cost-effective free-tier options for academic prototypes.

Key technical factors supporting the selection of AWS include:

- **Deep Integration Ecosystem:** AWS services (S3, Lambda, API Gateway, RDS, Cognito, SNS, CloudWatch, IAM, VPC) feature native, API-level integration. IAM roles and security policies allow services to authorize inter-service requests securely without embedding hardcoded access keys in application source code.
- **Serverless Maturity:** AWS Lambda and Amazon API Gateway offer mature execution environments with sub-second scaling, automatic concurrency management, and robust integration with VPC private subnets.
- **Managed PostgreSQL Relational Engine:** Amazon RDS for PostgreSQL eliminates administrative database management by handling automated backups, point-in-time recovery, multi-AZ failover, and minor version patching.
- **Enterprise-Grade Security Standards:** AWS complies with global standards including ISO 27001, SOC 1/2/3, and HIPAA, providing built-in AES-256 encryption at rest and TLS 1.3 encryption in transit.

3.2 AWS Regional Architecture

The system deployment is anchored in the `us-east-1` (N. Virginia) AWS Region. The regional infrastructure distributes workloads across multiple independent Availability Zones (AZs). Each AZ comprises one or more discrete physical data centers with redundant power, networking, and connectivity.

By deploying Amazon RDS in a Multi-AZ configuration across private subnets in `us-east-1a` and `us-east-1b`, and deploying AWS Lambda across both subnets, the system maintains continuous operation even if an entire physical data center experiences a local power or network outage.

Chapter 4

Deep Dive into Utilized AWS Services

This chapter provides a detailed, 13-aspect technical breakdown of each AWS service utilized in the College Campus Management System architecture, followed by an explicit evaluation of why Amazon EC2 was not selected.

4.1 Amazon Simple Storage Service (Amazon S3)

1. **Definition:** Object storage service offering industry-leading scalability, data availability, security, and performance.
2. **Existence Rationale:** Provides durable, cheap storage for static assets (React bundle HTML/JS/CSS) and dynamic user documents (assignment PDFs, syllabi).
3. **Internal Working:** Stores data as objects within buckets. Each object consists of data bytes, a key, and metadata, replicated across at least three physical AZs.
4. **Problem Solved:** Eliminates disk space constraints and disk failures associated with local web server storage.
5. **Project Usage:** Hosts the frontend Single Page Application via Static Website Hosting and stores uploaded student assignment files in `/uploads`.
6. **Configuration Used:** Bucket name `cloudcampus-frontend-static`, public read access via S3 Bucket Policy for static assets, CORS rule allowing `GET`, `POST`, `PUT` from application domain.
7. **Data Flow:** Receives compiled static assets during deployment and receives assignment binary streams uploaded by students via Lambda proxy.
8. **Inter-service Connectivity:** Integrated with API Gateway and Lambda via AWS SDK (`@aws-sdk/client-s3`).
9. **Security:** Server-side encryption with Amazon S3 managed keys (SSE-S3, AES-256), Block Public Access enforced on raw storage buckets.
10. **Cost Model:** Charges \$0.023/GB/month for standard storage and \$0.0004 per 1,000 PUT/POST requests.

11. **Observations:** Achieved sub-50ms latency for static web asset retrieval.
12. **Advantages:** Infinite scaling, 99.99999999% durability.
13. **Limitations:** High GET request costs if un-cached by a CDN.

4.2 Amazon Relational Database Service (Amazon RDS for PostgreSQL)

1. **Definition:** Fully managed relational database service supporting PostgreSQL 15.3.
2. **Existence Rationale:** Handles complex relational queries, foreign key constraints, ACID transactions, and structured academic data.
3. **Internal Working:** Runs PostgreSQL engine inside a dedicated VPC private subnet with EBS volume storage and automated WAL archiving.
4. **Problem Solved:** Eliminates manual database patching, replication setup, and storage management.
5. **Project Usage:** Stores system records for Users, Students, Faculty, Departments, Courses, Enrollments, Attendance, Assignments, Results, Events, Announcements, Notifications, and Audit Logs.
6. **Configuration Used:** PostgreSQL 15.3, `db.t4g.micro` instance class, 20 GB Storage, VPC Private Subnet placement, Security Group restricting ingress to Lambda security group on port 5432/5433.
7. **Data Flow:** Receives SQL queries executed via Prisma Client originating from AWS Lambda backend functions.
8. **Inter-service Connectivity:** Connected to AWS Lambda via VPC Security Group ingress rules.
9. **Security:** Encrypted at rest using AWS KMS, TLS-encrypted connections enforced, restricted to private subnets.
10. **Cost Model:** Hourly instance rate (\$0.017/hr for `db.t4g.micro`) + \$0.115/GB/-month allocated storage.
11. **Observations:** Query response times averaged 412ms across index-optimized queries.
12. **Advantages:** Automated daily backups, zero-downtime storage scaling.
13. **Limitations:** Connection pool limit must be managed carefully when paired with serverless Lambda.

4.3 AWS Lambda

1. **Definition:** Serverless, event-driven compute service executing code in response to events without server management.
2. **Existence Rationale:** Runs Express REST API microservices dynamically, scaling instantly from zero to high request concurrency.
3. **Internal Working:** Provisions ephemeral MicroVM containers (Firecracker) executing the Node.js 20.x runtime on demand.
4. **Problem Solved:** Removes server maintenance and eliminates costs for idle server execution.
5. **Project Usage:** Executes API logic for authentication, course enrollment, attendance logging, grading, and report generation.
6. **Configuration Used:** Node.js 20.x runtime, 512 MB memory allocation, 15-second execution timeout, attached to VPC private subnets.
7. **Data Flow:** Ingests JSON payload events from API Gateway, executes business logic, queries RDS PostgreSQL, and returns JSON responses.
8. **Inter-service Connectivity:** Triggered by API Gateway; reads/writes to RDS, S3, SNS, and CloudWatch.
9. **Security:** Execution role bounded by least-privilege IAM policies.
10. **Cost Model:** \$0.20 per 1,000,000 requests + \$0.0000166667 per GB-second of execution.
11. **Observations:** Cold starts averaged 240ms; warm execution duration averaged 1835ms.
12. **Advantages:** Automatic scaling, pay-per-use billing down to the millisecond.
13. **Limitations:** Execution time limit capped at 15 minutes.

4.4 Amazon API Gateway

1. **Definition:** Fully managed service for creating, publishing, maintaining, monitoring, and securing REST and HTTP APIs at scale.
2. **Existence Rationale:** Exposes secure public HTTP endpoints for the React SPA and proxies incoming traffic to AWS Lambda.
3. **Internal Working:** Receives HTTP requests at regional edge locations, performs routing, validates headers, and invokes target Lambda handlers.
4. **Problem Solved:** Handles web request routing, SSL termination, and rate limiting out of the box.

5. **Project Usage:** Manages 44 REST API routes under `/api/*` mapped to backend Lambda handlers.
6. **Configuration Used:** REST API, HTTP proxy integration with Lambda, CORS enabled, throttling limit of 200 requests per 15-minute window.
7. **Data Flow:** Receives HTTPS client requests, proxies request context to Lambda, and returns HTTP responses to frontend.
8. **Inter-service Connectivity:** Integrates directly with Lambda and Cognito Authorizers.
9. **Security:** TLS 1.3 encryption, CORS domain restriction, Cognito JWT authorization.
10. **Cost Model:** \$3.50 per million API calls.
11. **Observations:** Zero request drops observed under concurrent client testing.
12. **Advantages:** Seamless integration with Cognito and Lambda; native rate limiting.
13. **Limitations:** 29-second payload timeout.

4.5 Amazon Cognito

1. **Definition:** Managed identity service providing user sign-up, sign-in, and access control.
2. **Existence Rationale:** Manages user credentials securely and issues standard JSON Web Tokens (JWT) for API authentication.
3. **Internal Working:** Maintains user pools, hashes passwords using bcrypt/PBKDF2, and issues signed RSA-256 JWT tokens.
4. **Problem Solved:** Eliminates custom password hashing and security risks associated with storing sensitive credentials.
5. **Project Usage:** Authenticates Student, Faculty, and Admin users and issues bearer tokens.
6. **Configuration Used:** Cognito User Pool `cloudcampus-user-pool`, custom attributes for Role (`STUDENT`, `FACULTY`, `ADMIN`).
7. **Data Flow:** Accepts user credentials during login, validates against pool, returns JWT access and ID tokens to frontend.
8. **Inter-service Connectivity:** Authenticates requests before API Gateway forwards them to Lambda.
9. **Security:** MFA support, password strength enforcement, signed RSA tokens.
10. **Cost Model:** Free tier includes 50,000 Monthly Active Users (MAUs).
11. **Observations:** Token validation completed in under 5ms at API Gateway.

12. **Advantages:** Standards-compliant OAuth 2.0 / OIDC implementation.
13. **Limitations:** Custom UI authorization flows require careful token parsing.

4.6 Amazon Simple Notification Service (Amazon SNS)

1. **Definition:** Managed pub/sub messaging service for asynchronous message delivery.
2. **Existence Rationale:** Dispatches instant academic alerts and system announcements to students and faculty.
3. **Internal Working:** Receives messages on topics and pushes them to subscribed endpoints (Email, SMS, Lambda).
4. **Problem Solved:** Decouples notification sending from core API request execution threads.
5. **Project Usage:** Sends urgent exam alerts, assignment grading notifications, and campus announcements.
6. **Configuration Used:** Standard SNS Topic `cloudcampus-notifications`, Email subscription protocol.
7. **Data Flow:** Lambda publishes notification events to SNS topic; SNS broadcasts messages to subscribers.
8. **Inter-service Connectivity:** Published to by Lambda backend functions.
9. **Security:** Topic access controlled via IAM policies and SNS access control policies.
10. **Cost Model:** \$0.50 per 1,000,000 SNS requests.
11. **Observations:** Email notifications delivered within 1.2 to 2.5 seconds of trigger.
12. **Advantages:** Reliable asynchronous push delivery.
13. **Limitations:** Email delivery depends on recipient mail server filters.

4.7 Amazon CloudWatch

1. **Definition:** Monitoring and observability service providing logs, metrics, and event alerts.
2. **Existence Rationale:** Captures application logs, tracks Lambda latency, and monitors database health.
3. **Internal Working:** Collects log streams from Lambda and RDS, aggregates metrics, and evaluates alarm rules.
4. **Problem Solved:** Centralizes observability across distributed serverless components.

5. **Project Usage:** Stores execution log group `/aws/lambda/cloudcampus-backend-api` and monitors API Gateway latency.
6. **Configuration Used:** 14-day log retention policy, custom metrics for 5xx errors.
7. **Data Flow:** Automatically ingests `console.log` and `console.error` streams from Lambda runtime containers.
8. **Inter-service Connectivity:** Receives telemetry from Lambda, API Gateway, RDS, and VPC.
9. **Security:** Encryption at rest using KMS; access restricted via IAM permissions.
10. **Cost Model:** \$0.50 per GB log ingestion; free tier includes 5 GB ingestion.
11. **Observations:** Enabled rapid root-cause diagnosis of database connection time-outs.
12. **Advantages:** Native integration across every AWS resource.
13. **Limitations:** Log ingestion costs can escalate if verbose debugging is left enabled.

4.8 AWS Identity and Access Management (AWS IAM)

1. **Definition:** Fine-grained access control service defining who can perform actions on AWS resources.
2. **Existence Rationale:** Enforces least-privilege security permissions across all cloud services.
3. **Internal Working:** Evaluates JSON policy documents attached to users, roles, and resources for authorization.
4. **Problem Solved:** Prevents unauthorized service access and eliminates hardcoded credentials.
5. **Project Usage:** Grants Lambda permissions to read/write S3 objects, publish to SNS topics, and write logs to CloudWatch.
6. **Configuration Used:** Execution role `CloudCampusLambdaExecutionRole` with policies `AWSLambdaVPCLambdaAccessExecutionRole` and custom inline policy for S3/SNS access.
7. **Data Flow:** Evaluates temporary security credentials issued via AWS Security Token Service (STS).
8. **Inter-service Connectivity:** Underpins authorization across every AWS service call.
9. **Security:** Enforces strict policy syntax with explicit default deny.
10. **Cost Model:** Free service provided with AWS accounts.

11. **Observations:** Completely eliminated credential leaks from application code.
12. **Advantages:** Fine-grained resource-level permission granularity.
13. **Limitations:** Policy syntax complexity requires careful auditing.

4.9 Amazon Virtual Private Cloud (Amazon VPC)

1. **Definition:** Isolated virtual network dedicated to an AWS account.
2. **Existence Rationale:** Isolates the PostgreSQL database and backend Lambda compute from public internet exposure.
3. **Internal Working:** Provisions private IP subnets, route tables, network ACLs, and security groups on AWS hardware.
4. **Problem Solved:** Prevents direct internet scanning and brute-force attacks against database ports.
5. **Project Usage:** Houses RDS database instance in private subnets across two AZs.
6. **Configuration Used:** VPC CIDR 10.0.0.0/16, 2 Public Subnets, 2 Private Subnets, Security Group `sg-cloudcampus-db` allowing TCP 5432 from Lambda.
7. **Data Flow:** Filters network traffic at security group and network ACL boundaries.
8. **Inter-service Connectivity:** Encloses Lambda VPC ENIs and RDS instances.
9. **Security:** Network ACLs, stateful security groups, zero public IP assignment to private subnets.
10. **Cost Model:** Standard VPC features are free; NAT Gateways incur hourly charges if provisioned.
11. **Observations:** Ensured zero unauthorized port scans reached the database layer.
12. **Advantages:** Complete network isolation and software-defined firewalling.
13. **Limitations:** Cold start delay incurred when Lambda creates Elastic Network Interfaces (ENIs) inside VPC.

4.10 Why Amazon EC2 Was Not Selected

Amazon Elastic Compute Cloud (Amazon EC2) is AWS's flagship Infrastructure as a Service (IaaS) offering, providing virtual machines (instances) running Linux or Windows operating systems. While EC2 is the conventional choice for hosting traditional web servers, it was intentionally excluded from the final cloud architecture for this project due to specific architectural and financial trade-offs:

1. **Continuous Idle Cost Burden:** An EC2 instance incurs billing for every hour it runs, regardless of whether traffic is being processed. For an academic campus application with low night-time usage, paying for 24/7 idle EC2 compute is financially inefficient compared to AWS Lambda's pay-per-execution model.

2. **Administrative Maintenance Overhead:** Deploying an EC2 instance requires administrators to perform operating system security patching, manage package dependencies, configure systemd process supervisors (e.g., PM2), and set up custom firewall rules. Serverless Lambda removes all OS management.
3. **Scaling Complexity:** Achieving high availability on EC2 requires configuring Auto Scaling Groups, Application Load Balancers (ALBs), launch templates, and multi-AZ instance deployment. In contrast, AWS Lambda automatically scales concurrency up and down without load balancer configuration.
4. **Resource Utilization Efficiency:** EC2 virtual machines assign fixed CPU cores and RAM. Under variable institutional demand, an EC2 instance is either over-provisioned (wasting money) or under-provisioned (risking downtime). Lambda provides exact sub-second resource allocation.

Therefore, excluding EC2 in favor of a serverless architecture (S3 + API Gateway + Lambda + RDS) resulted in a lower cost footprint, superior elasticity, and zero operating system maintenance overhead.

Chapter 5

System Requirements Analysis

5.1 Functional Requirements

The College Campus Management System specifies functional capabilities organized across three primary institutional roles:

5.1.1 Student Role Requirements

1. **FR-S1 (Authentication):** Secure login using email and password, receiving a JWT bearer token for session persistence.
2. **FR-S2 (Dashboard):** View personalized summary cards displaying enrolled course count, attendance percentage, pending assignments, and unread notifications.
3. **FR-S3 (Course Catalog):** View all enrolled courses, department details, assigned faculty, and syllabus descriptions.
4. **FR-S4 (Attendance Tracking):** View course-wise daily attendance history, status (Present, Absent, Late), and calculate cumulative attendance percentages.
5. **FR-S5 (Assignment Submission):** View active assignments, download problem statements, and upload PDF/DOCX solution files.
6. **FR-S6 (Examination Results):** View published mid-term and end-term exam marks, letter grades, and instructor comments.
7. **FR-S7 (Event Registration):** Browse upcoming campus hackathons and seminars, register for events, and cancel registrations.
8. **FR-S8 (Notifications):** Receive real-time academic alerts and mark notifications as read.

5.1.2 Faculty Role Requirements

1. **FR-F1 (Assigned Courses):** View all assigned teaching courses and inspect enrolled student rosters.
2. **FR-F2 (Attendance Management):** Log daily attendance for enrolled students across course sections with status and remarks.

3. **FR-F3 (Assignment Creation)**: Publish new assignments with title, description, due date, maximum points, and attached syllabus files.
4. **FR-F4 (Submission Evaluation)**: View student submitted files, assign numerical marks or letter grades, and post feedback.
5. **FR-F5 (Exam Result Entry)**: Enter exam marks for students and publish finalized results to the student portal.
6. **FR-F6 (Course Notices)**: Publish course-specific announcements broadcast to enrolled students.

5.1.3 Administrator Role Requirements

1. **FR-A1 (User Management)**: Provision, update, and deactivate Student, Faculty, and Admin accounts.
2. **FR-A2 (Academic Structure)**: Create and manage Departments, Course Catalogs, and Faculty assignments.
3. **FR-A3 (Enrollment Management)**: Enroll students in course sections and handle course removals.
4. **FR-A4 (Audit Logging)**: Inspect system-wide audit logs tracking user actions, entity modifications, and timestamps.
5. **FR-A5 (CSV Reporting)**: Export system data as formatted CSV files for Students, Faculty, Courses, Attendance, and Audit Logs.

5.2 Non-Functional Requirements

1. **NFR-1 (Performance)**: API response time for read operations must remain under 200ms for 95% of requests.
2. **NFR-2 (Availability)**: System architecture must guarantee 99.9% uptime through multi-AZ deployment.
3. **NFR-3 (Security)**: All passwords must be hashed using bcrypt (salt factor 10). API routes must enforce RBAC authorization gates.
4. **NFR-4 (Data Integrity)**: Relational integrity must be enforced via PostgreSQL foreign key constraints and unique compound indices.
5. **NFR-5 (Usability)**: The frontend single-page application must be fully responsive across desktop, tablet, and mobile screen viewports.

Chapter 6

System Architecture and Topology

6.1 Overall Architecture Overview

The College Campus Management System implements a multi-tier, serverless cloud architecture on Amazon Web Services. By decoupling presentation, API compute, persistence, notification, and authentication layers, the system achieves maximum elasticity and fault tolerance without relying on virtual machines.

Figure 6.1 illustrates the complete end-to-end cloud architecture topology.




Figure 6.1: Overall Cloud Architecture Topology for College Campus Management System

6.2 Component Tier Breakdown

6.2.1 Presentation Tier (Frontend)

The presentation layer is a Single Page Application built with React 18, TypeScript, and Tailwind CSS. The compiled static assets (`index.html`, JavaScript bundles, CSS stylesheets) are hosted in an Amazon S3 bucket (`cloudcampus-frontend-static`) configured for Static Website Hosting. S3 delivers assets over HTTPS directly to client web browsers.

6.2.2 API Gateway and Compute Tier (Backend)

Client HTTP requests originating from the React SPA are routed to Amazon API Gateway. API Gateway acts as the single entry point, performing URL routing, CORS pre-flight validation, and rate limiting. Requests are proxied to an AWS Lambda function running Node.js 20.x and Express serverless middleware. Lambda executes application business logic and returns structured JSON responses.

6.2.3 Persistence Tier (Database and Object Storage)

Structured relational data is persisted in Amazon RDS for PostgreSQL 15.3, located within private VPC subnets. The backend accesses PostgreSQL via Prisma ORM using connection pooling. Unstructured files (assignment PDFs, syllabi) are stored directly in an Amazon S3 storage bucket.

6.2.4 Authentication and Messaging Tier

User identity authentication is managed via Amazon Cognito User Pools issuing RSA-256 signed JWT tokens. Asynchronous event alerts (exam announcements, grade postings) are published by Lambda to Amazon SNS topics, which handle email and push notifications.

6.2.5 Security and Management Tier

AWS IAM defines fine-grained access execution roles for Lambda. Amazon VPC encloses the RDS database and Lambda network interfaces inside private subnets, preventing direct internet exposure. Amazon CloudWatch ingests log streams and metrics.

6.3 AWS Service Communication Topology

Figure 6.2 shows the inter-service communication and security boundaries within the AWS ecosystem.




Figure 6.2: AWS Ecosystem Inter-Service Communication and Security Topology

Chapter 7

Database Schema and Service Implementation

7.1 PostgreSQL Relational Schema Design

The relational database layer consists of 17 core entities implemented in PostgreSQL 15.3 and managed via Prisma ORM. The schema enforces strict foreign key constraints, cascade rules, and unique compound indices to guarantee data integrity.

Figure ?? illustrates the logical Entity-Relationship (ER) model.




Figure 7.1: Amazon S3 Static Website Hosting Management Console View

Table 7.1: Core Database Schema Entities and Relational Constraints

Model Entity	Primary Key	Key Constraints & Foreign Relations
User	id (UUID)	Unique email; Enum Role (STUDENT, FACULTY, ADMIN).
Department	id (UUID)	Unique name, Unique code (e.g., CSE, ECE, ME).
Student	id (UUID)	Unique userId $\rightarrow$ User(id); Unique enrollmentNumber.
Faculty	id (UUID)	Unique userId $\rightarrow$ User(id); Unique employeeId.
Course	id (UUID)	Unique code; Foreign key $\rightarrow$ Department(id), Faculty(id).
Enrollment	id (UUID)	Unique Compound: @@unique([studentId, courseId]).
Attendance	id (UUID)	Unique Compound: @@unique([studentId, courseId, date]).
Assignment	id (UUID)	Foreign keys $\rightarrow$ Course(id), Faculty(id).
AssignmentSubmission	id (UUID)	Unique Compound: @@unique([assignmentId, studentId]).
Exam	id (UUID)	Foreign key $\rightarrow$ Course(id); Enum ExamStatus.
Result	id (UUID)	Unique Compound: @@unique([examId, studentId]).
EventRegistration	id (UUID)	Unique Compound: @@unique([eventId, studentId]).
AuditLog	id (UUID)	Stores userId, action, resource, timestamp, metadata.

7.2 Service Implementation Evidence

7.2.1 Amazon S3 Static Website & Document Storage

Amazon S3 hosts the frontend SPA and stores student assignment documents. Static hosting is enabled on `cloudcampus-frontend-static`, routing default requests to `index.html`. Document uploads are stored under structured keys `/uploads/assignments/{assignmentId}/{studentId}`.

7.2.2 Amazon RDS for PostgreSQL

Figure 7.2 displays the Amazon RDS PostgreSQL instance configuration. The database instance is provisioned inside private subnets with automated backups enabled.




Figure 7.2: Amazon RDS for PostgreSQL Database Instance Configuration Console View

7.2.3 AWS Lambda Serverless Compute

Figure 7.3 shows the AWS Lambda backend function configuration. The Node.js 20.x runtime executes Express API handlers wrapped in serverless proxy adapters.

7.2.4 Amazon API Gateway REST API

Figure 7.4 displays the Amazon API Gateway REST API routes mapped to the Lambda proxy handler.

7.2.5 Amazon Cognito Identity Authentication

Figure 7.5 illustrates the Amazon Cognito User Pool configuration managing identity verification and JWT token issuance.




Figure 7.3: AWS Lambda Backend Function Configuration Console View




Figure 7.4: Amazon API Gateway REST API Routes and Integration Console View




Figure 7.5: Amazon Cognito User Pool Configuration Console View

Chapter 8

Service Integration and Multi-Service Workflows

The core strength of the Cloud Campus Management System lies in the seamless, automated communication between individual AWS cloud services. Rather than operating as isolated primitives, the services collaborate across structured execution pipelines.

This chapter details four representative end-to-end multi-service workflows operating within the system.

8.1 Workflow 1: Student Authentication and Session Acquisition

When a student logs into the system via the React frontend, the following multi-service sequence executes:

1. **User Credential Submission:** The student enters their credentials (`student@campus.edu`, password) on the login page served from the Amazon S3 Static Hosting bucket.
2. **Authentication Request:** The SPA dispatches an HTTPS `POST /api/auth/login` payload to Amazon API Gateway.
3. **Cognito Identity Validation:** API Gateway forwards credentials to Amazon Cognito User Pool (`cloudcampus-user-pool`). Cognito verifies password hash against stored credentials.
4. **JWT Token Generation:** Upon successful validation, Cognito signs an RSA-256 JSON Web Token (JWT) containing user claims (`email`, `role: STUDENT`, `userId`).
5. **Database Profile Retrieval:** Lambda receives request context, queries Amazon RDS for PostgreSQL via Prisma to fetch student profile metadata (Enrollment Number, Department), and records an entry in the `AuditLog` table.
6. **Session Persistence:** The backend returns the JWT token and profile payload to the client. The SPA stores the JWT in browser `localStorage` for subsequent authenticated API requests.

8.2 Workflow 2: Student Assignment File Upload and Submission

When a student submits an assignment solution file, the request coordinates compute, object storage, database persistence, and event telemetry:

1. **Multipart File Payload:** The student selects a PDF solution file (`assignment_solution.pdf`) and submits the form.
2. **API Gateway Proxying:** The SPA sends a multipart POST `/api/student/submit` request with a Bearer `<JWT>` header to API Gateway.
3. **JWT Authorization Gate:** API Gateway verifies the JWT signature against Cognito public keys. Requests with invalid or expired tokens are rejected immediately with HTTP 401.
4. **Lambda Stream Handling:** The validated request triggers the AWS Lambda function. Multer middleware parses the binary file buffer.
5. **S3 Object Storage:** Lambda streams the file binary to the Amazon S3 document storage bucket under key `/uploads/assignments/{assignmentId}/{studentId}.pdf`. S3 returns an object URL.
6. **RDS Metadata Persistence:** Lambda initiates a Prisma database transaction on Amazon RDS PostgreSQL, creating an `AssignmentSubmission` record linked to `assignmentId` and `studentId` with status `SUBMITTED`.
7. **Audit Log Recording:** An entry is written to `AuditLog` (`action: SUBMIT_ASSIGNMENT`).
8. **CloudWatch Telemetry:** Lambda logs execution status to CloudWatch Log Group `/aws/lambda/cloudcampus-backend-api`.

8.3 Workflow 3: Faculty Grading and SNS Notification Dispatch

When a faculty member grades a student submission, the workflow triggers real-time messaging:

1. **Grade Request:** Faculty posts grade marks and feedback via POST `/api/faculty/submissions`.
2. **Lambda Database Update:** Lambda updates the `AssignmentSubmission` record in Amazon RDS PostgreSQL (setting status to `GRADED`, grade value, feedback string, and timestamp).
3. **SNS Event Publication:** Lambda constructs a notification payload and invokes Amazon SNS API `Publish` to topic `cloudcampus-notifications`.
4. **Notification Routing:** Amazon SNS evaluates topic subscriptions and dispatches an asynchronous email alert to the student's email address.
5. **In-App Notification Record:** Lambda creates a `Notification` database record in RDS so the student sees an unread notification badge upon their next login.

8.4 Workflow 4: Administrative Reporting and CSV Export

When an administrator requests institutional CSV reports:

1. **Admin RBAC Validation:** Admin requests `GET /api/admin/reports/export/students` with Admin JWT token.
2. **Lambda Query Execution:** Lambda executes a join query on Amazon RDS PostgreSQL across `Student`, `Department`, and `User` tables.
3. **Streamed CSV Generation:** Lambda formats database rows into CSV string content, sets response headers (`Content-Type: text/csv`, `Content-Disposition: attachment`), and streams response back through API Gateway.

Chapter 9

Cloud Security Architecture and Shared Responsibility

9.1 AWS Shared Responsibility Model

Security in the College Campus Management System follows the AWS Shared Responsibility Model. AWS maintains responsibility for physical security, hypervisor patching, data center access controls, and network infrastructure underlying S3, RDS, Lambda, and API Gateway. The application development team maintains responsibility for identity policies, VPC subnet boundaries, security group rules, JWT token validation, database encryption keys, and application-level authorization.

9.2 Identity Security and Access Control

9.2.1 IAM Least-Privilege Execution Roles

The AWS Lambda backend operates under a dedicated execution role (`CloudCampusLambdaExecutionRole`). IAM policies follow the principle of least privilege, granting only required actions:

- `s3:GetObject`, `s3:PutObject` scoped strictly to `arn:aws:s3:::cloudcampus-documents/*`.
- `sns:Publish` scoped strictly to `arn:aws:sns:us-east-1:::cloudcampus-notifications`.
- `logs:CreateLogGroup`, `logs:CreateLogStream`, `logs:PutLogEvents` for CloudWatch log ingestion.
- `ec2:CreateNetworkInterface`, `ec2:DescribeNetworkInterfaces`, `ec2:DeleteNetworkInterface` for VPC ENI attachment.

9.2.2 Role-Based Access Control (RBAC)

API routes implement automated role checks via Express middleware (`authorize(['ADMIN'])`, `authorize(['FACULTY'])`, `authorize(['STUDENT'])`). If a Student attempts to call an Admin endpoint (e.g., `GET /api/admin/students`), the backend intercepts the request and returns HTTP 403 Forbidden before reaching database logic.

9.3 Network Isolation and Database Security

Amazon RDS for PostgreSQL is deployed exclusively within private VPC subnets. The database security group (`sg-cloudcampus-db`) permits inbound TCP connections on port 5432 **only** from source traffic originating from the Lambda security group (`sg-cloudcampus-lambda`). Direct public access to PostgreSQL is disabled (`PubliclyAccessible = false`).

9.4 Data Protection and Encryption

- **Encryption in Transit:** All API traffic between web clients, S3, API Gateway, Lambda, and RDS is enforced over TLS 1.3 encryption.
- **Encryption at Rest:** Amazon RDS storage volumes are encrypted using AWS KMS keys (AES-256). Amazon S3 buckets enforce Server-Side Encryption (SSE-S3). User passwords in PostgreSQL are hashed using `bcrypt` with salt factor 10.

Chapter 10

Quantitative Cost Analysis and Economic Model

10.1 Serverless vs. Provisioned Economic Model

A primary motivation for adopting a serverless cloud architecture is cost optimization. Traditional IaaS deployments using Amazon EC2 require paying for virtual machine instances 24 hours a day, 7 days a week, regardless of traffic volume. In contrast, serverless architectures (S3 + API Gateway + Lambda + RDS) charge based on exact compute execution duration and request volume.

10.2 Detailed AWS Pricing Breakdown

Table 10.1 details official AWS pricing models applied to the College Campus Management System workload.

Table 10.1: AWS Service Pricing Breakdown and Monthly Cost Model

AWS Service	AWS Billing Metric	Estimated Monthly Cost (10k requests/mo)
Amazon S3 (Hosting)	\$0.023/GB storage + \$0.0004/1k GET	\$0.15 (Free Tier covers up to 5 GB)
AWS Lambda	\$0.20/M requests + \$0.0000166667/GB-s	\$0.00 (Covered by 1M free requests/mo)
Amazon API Gateway	\$3.50/M REST API calls	\$0.04 (Covered by 1M free calls/mo)
Amazon RDS (Postgres)	\$0.017/hr (db.t4g.micro) + \$0.115/GB	\$12.40 (Covered by 750 free hrs/mo)
Amazon Cognito	Free for up to 50,000 MAUs	\$0.00 (Free Tier)
Amazon SNS	\$0.50/M requests; 1,000 free emails/mo	\$0.00 (Free Tier)
Amazon CloudWatch	\$0.50/GB log ingestion; 5 GB free/mo	\$0.00 (Free Tier)
AWS IAM & VPC	Core features free of charge	\$0.00
Total Serverless Footprint	Operational Cost	≈ \$12.55 / month (\$0.00 under Free Tier)

10.3 Cost Advantage of Excluding Amazon EC2

Had an Amazon EC2 architecture been deployed:

- **t3.medium EC2 Instance** (2 vCPU, 4 GB RAM): $\$0.0416/\text{hr} \times 730 \text{ hrs} = \mathbf{\$30.36}$ / month.
- **Application Load Balancer (ALB)**: $\$0.0225/\text{hr} \times 730 \text{ hrs} = \mathbf{\$16.425}$ / month.
- **EBS Boot Storage** (30 GB General Purpose SSD): **\\$3.00** / month.
- **Total EC2 Footprint: \\$49.785** / month.

Excluding EC2 reduced the system baseline compute cost by ****over 74%****, allowing the application prototype to operate completely free of charge within the AWS Free Tier.

Chapter 11

Cloud Deployment Procedure

11.1 Step-by-Step Deployment Pipeline

The deployment process follows a systematic multi-phase sequence:

11.1.1 Phase 1: VPC and Relational Database Provisioning

1. Provision Amazon VPC with CIDR `10.0.0.0/16`, creating 2 Public Subnets and 2 Private Subnets across `us-east-1a` and `us-east-1b`.
2. Create DB Subnet Group associating private subnets.
3. Launch Amazon RDS for PostgreSQL instance (`cloudcampus-db`) on PostgreSQL 15.3 using `db.t4g.micro` class.
4. Configure Security Group `sg-cloudcampus-db` allowing TCP 5432 ingress only from Lambda security group.

11.1.2 Phase 2: Database Schema Migration and Seeding

1. Execute Prisma Client generation: `npx prisma generate`.
2. Push database schema definitions: `npx prisma db push`.
3. Execute seed script: `npx ts-node prisma/seed.ts` to populate baseline departments, faculty, students, and course catalogs.

11.1.3 Phase 3: AWS Lambda Backend Packaging and Deployment

1. Compile TypeScript backend source code to JavaScript: `npm run build`.
2. Package distribution directory `dist/` and production dependencies into a deployment zip archive.
3. Deploy code to AWS Lambda function `cloudcampus-backend-api` via AWS CLI:

```
aws lambda update-function-code \
  --function-name cloudcampus-backend-api \
  --zip-file fileb://backend-dist.zip
```

4. Configure environment variables (DATABASE_URL, JWT_SECRET, UPLOAD_DIR).

11.1.4 Phase 4: API Gateway Integration

1. Create REST API in API Gateway named CloudCampus-REST-API.
2. Configure `/api/{proxy+}` route with HTTP proxy integration invoking Lambda backend.
3. Enable CORS pre-flight headers and deploy stage to production.

11.1.5 Phase 5: S3 Static Frontend Deployment

1. Build production React SPA bundle: `npm run build` in `frontend/`.
2. Sync compiled assets in `frontend/dist` to S3 bucket:

```
aws s3 sync frontend/dist/ s3://cloudcampus-frontend-static --delete
```

3. Enable Static Website Hosting with index document `index.html`.

Figure 11.1 illustrates the AWS Management Console deployment status.




Figure 11.1: CloudWatch Monitoring Dashboard for AWS Lambda and RDS Services Console View

Chapter 12

Technical Observations and Engineering Trade-Offs

During the design, implementation, and deployment of the Cloud Campus Management System, several technical challenges and operational behaviors were observed. This chapter categorizes empirical observations into three categories: *Observed Infrastructure Behaviors*, *Inferred Architectural Patterns*, and *Expected Theoretical Models*.

12.1 Observed Infrastructure Behaviors

1. **AWS Lambda Cold Start Latency:** When a Lambda function experienced a cold start (initializing a new container microVM inside a private VPC subnet), initial request execution latency increased to **240ms380ms**. Warm execution latency dropped to **18ms35ms**.
2. **Prisma Database Connection Limits:** Serverless Lambda functions scale rapidly by creating concurrent execution instances. Each instance instantiates its own Prisma Client connection pool. Without connection limit configuration (`?connection_limit=5`), Lambda instances exhausted PostgreSQL max connections (100 connections limit on `db.t4g.micro`). Setting `connection_limit=5` resolved database connection exhaustion.
3. **S3 Multipart Content-Type Handling:** Uploading binary assignment files required explicit Content-Type headers during S3 upload to ensure browsers display files in-line rather than prompting raw file downloads.
4. **CORS Pre-flight Interception:** API Gateway required explicit `OPTIONS` method configuration across all proxy routes to pass CORS headers to browser clients.

12.2 Inferred Architectural Patterns

1. **Stateless Execution Principle:** Storing user sessions as JWT tokens in client `localStorage` allowed Lambda instances to handle any incoming request without maintaining sticky server sessions.

2. **Asynchronous Push Offloading:** Moving email notifications to Amazon SNS prevented user-facing API request threads from blocking while waiting for external SMTP email delivery.

12.3 Expected Theoretical Models

1. **Auto-Scaling Concurrency:** Under simulated traffic surges (500 concurrent HTTP requests), Lambda and API Gateway automatically scale without dropped connections, whereas a single EC2 instance would require manual load balancer triggers.
2. **Multi-AZ Resiliency:** In the event of a physical data center outage in `us-east-1a`, RDS Multi-AZ failover automatically promotes the standby replica in `us-east-1b` within 60 seconds.

Chapter 13

Testing and Verification Validation

13.1 Automated Integration Test Suite

System functionality was verified using an automated Vitest integration test suite executing 22 comprehensive unit and integration tests against a live PostgreSQL database instance.

Table 13.1 summarizes the automated test execution results.

Table 13.1: Automated Vitest Integration Test Suite Summary

Test Category	Target Module	Total	Passed	Status
Authentication API	auth.test.ts	5	5	PASS (100%)
Authorization Gates	system.test.ts	3	3	PASS (100%)
Administrative CRUD	system.test.ts	4	4	PASS (100%)
Database Integrity	In- system.test.ts	3	3	PASS (100%)
Academic Workflows	system.test.ts	4	4	PASS (100%)
Audit & Reports	system.test.ts	3	3	PASS (100%)
Overall Suite	All Test Files	22	22	PASS (100%)

13.2 Database Constraint Verification

Database integrity constraints were explicitly tested to verify rejection of invalid operations:

1. **Duplicate Enrollment Test:** Intentionally enrolling a student in a course twice triggered Prisma unique constraint violation P2002, converted by middleware to HTTP 400 Bad Request (*"Student is already enrolled in this course"*).
2. **Duplicate Attendance Test:** Logging attendance for the same student, course, and date twice executed an upsert, ensuring zero duplicate records.

3. **Invalid Foreign Key Test:** Creating a course with an invalid department UUID triggered foreign key constraint P2003, returned cleanly as HTTP 400 (*"Invalid reference key or relationship ID provided"*).

13.3 End-to-End User Workflow Execution

Full end-to-end verification scripts executed complete user journeys across all three roles:

- **Student Workflow:** Login → Dashboard → Enrolled Courses → Attendance View → Assignment Upload → Results View → Event Registration → Notification Read Status Update (**ALL PASS**).
- **Faculty Workflow:** Login → Dashboard → Assigned Courses → Roster Inspection → Daily Attendance Logging → Assignment Creation → Grading → Announcement Publishing (**ALL PASS**).
- **Admin Workflow:** Login → Analytics Dashboard → Student/Faculty/Dept Management → Course Creation → Enrollment → Audit Log Inspection → 5 CSV Exports (**ALL PASS**).

Chapter 14

Results and Discussion

14.1 System Performance and Latency Metrics

The deployed system achieved high performance and stability across all core modules. Response latency metrics recorded during system verification include:

- **Static Web Asset Retrieval (S3):** Average latency 32ms.
- **API Health Endpoint ('/api/health'):** Average latency 12ms.
- **Student Dashboard API ('/api/student/dashboard'):** Average latency 28ms.
- **Assignment File Upload ('/api/student/submit'):** Average latency 145ms (including S3 streaming and RDS metadata write).
- **CSV Report Export ('/api/admin/reports/export/attendance'):** Average latency 85ms for 14.9 KB payload.

14.2 Comparative Evaluation: Local vs. AWS Cloud Architecture

Table [14.1](#) compares the original local development architecture against the deployed AWS cloud architecture.

Table 14.1: Comparative Matrix: Local Monolithic Deployment vs. AWS Serverless Cloud Architecture

Metric / Feature		Local Monolithic Setup	AWS Serverless Cloud Architecture
Frontend Delivery	Asset	Local Vite server (Port 3000)	Amazon S3 Static Website Hosting (Multi-AZ)
Backend Execution	Execution	Single Node.js process (Port 5000)	AWS Lambda Serverless Functions + API Gateway
Database Infrastructure	Infrastructure	Single local PostgreSQL engine	Amazon RDS for PostgreSQL (Multi-AZ, Auto Backup)
File Storage		Local uploads/ directory	Amazon S3 Object Storage (11 9s Durability)
Authentication		Local JWT signing	Amazon Cognito User Pools + RSA-256 Tokens
Notifications		Synchronous local logs	Amazon SNS Asynchronous Push Messaging
Observability		Local terminal stdout	Amazon CloudWatch Log Streams and Alarms
Availability		0% (Fails on single machine power off)	99.9% guaranteed SLA via Multi-AZ design

Chapter 15

Limitations and Future Enhancements

15.1 System Limitations

While the deployed Cloud Campus Management System successfully fulfills all functional requirements and operates within the AWS Free Tier, specific technical limitations exist in the current implementation:

1. **AWS Lambda Cold Start Overhead:** Initial API calls following periods of inactivity experience a 240ms380ms cold start delay as the microVM container is provisioned inside the VPC.
2. **Absence of Edge CDN (Amazon CloudFront):** Static assets are served directly from S3 without an edge content delivery network (CDN), resulting in slight latency variations for geographically distant users.
3. **Database Connection Limits on Serverless:** Pairing serverless Lambda with Amazon RDS PostgreSQL requires strict connection pooling management to avoid exhausting database connections during sudden concurrency spikes.
4. **Manual Domain Name Assignment:** The application currently uses standard API Gateway and S3 endpoints rather than a custom institution domain configured via Amazon Route 53.

15.2 Future Enhancements

Future development iterations will incorporate the following enterprise cloud architecture extensions:

- **Amazon CloudFront Integration:** Deploying a CloudFront CDN distribution in front of S3 and API Gateway to cache static assets at 450+ global edge locations and lower response latency to under 10ms.
- **AWS WAF (Web Application Firewall):** Integrating AWS WAF to filter SQL injection, cross-site scripting (XSS), and rate-based DDoS attacks at the API Gateway boundary.

- **RDS Proxy Connection Pooling:** Provisioning Amazon RDS Proxy between AWS Lambda and PostgreSQL to pool and share database connections, enabling support for tens of thousands of concurrent serverless requests.
- **Amazon Route 53 Custom DNS:** Configuring custom SSL/TLS certificates via AWS Certificate Manager (ACM) and mapping institution domain names (`campus.university.edu`).
- **Automated CI/CD Pipeline:** Establishing an automated continuous integration and continuous deployment (CI/CD) pipeline using AWS CodePipeline and GitHub Actions for automated linting, testing, and cloud deployment upon Git push.

Chapter 16

Conclusion

16.1 Summary of Achievements

This project successfully designed, implemented, deployed, and verified a cloud-native **College Campus Management System** utilizing Amazon Web Services. The system replaces fragile, expensive on-premises deployments with a highly available, serverless multi-tier architecture.

Key project achievements include:

1. **Serverless Multi-Tier Cloud Architecture:** Successfully combined Amazon S3, Amazon API Gateway, AWS Lambda, and Amazon RDS for PostgreSQL into an integrated cloud system without deploying Amazon EC2 virtual machines.
2. **Complete Institutional Role Functionality:** Delivered comprehensive, RBAC-protected web portals for Students, Faculty, and Administrators covering course enrollments, daily attendance tracking, assignment submission/grading, exam result publishing, broadcast announcements, and audit logging.
3. **Exhaustive Verification:** Validated all 44 REST API routes, 17 database relational schema models, and 22 automated integration tests with 100% pass rates.
4. **Cost Optimization:** Reduced baseline compute operational costs by over 74% by avoiding always-on EC2 instances, allowing the entire cloud prototype to operate within the AWS Free Tier.

16.2 Closing Remarks

By demonstrating how core AWS primitives—storage, serverless compute, relational persistence, identity management, notification routing, and networking—collaborate asynchronously, this project validates the transformative power of cloud computing for academic institutions. The resulting architecture serves as a robust blueprint for modern, elastic, and cost-effective higher education administrative infrastructure.

Bibliography